

```
In [1]: import zipfile
import os
zip_path = '/kaggle/input/dogs-vs-cats/train.zip'
extract_path = '/kaggle/working/'

with zipfile.ZipFile(zip_path, 'r') as zip_ref:
    zip_ref.extractall(extract_path)

dir_train = os.path.join(extract_path, 'train')

In [2]: import cv2
import numpy as np
from skimage.feature import hog
from skimage.color import rgb2gray
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, classification_report

In [3]: image_files = [os.path.join(dir_train, file) for file in os.listdir(dir_train) if file.endswith(('.jpg'))]

In [4]: imagefiles1 = image_files[:25000]

In [5]: def compute_hog_features(image):
image = rgb2gray(image)
features = hog(image, orientations=9, pixels_per_cell=(8, 8),
               cells_per_block=(2, 2), block_norm='L2-Hys',
               transform_sqrt=True)
return features

In [6]: def compute_hog_feature(image):

    gray_image = rgb2gray(image)

    features, hog_image = hog(gray_image, orientations=9, pixels_per_cell=(8, 8),
                              cells_per_block=(2, 2), block_norm='L2-Hys',
                              visualize=True, transform_sqrt=True)

    hog_image_rescaled = exposure.rescale_intensity(hog_image, in_range=(0, 20))

    return features, hog_image_rescaled

In [7]: total_images = len(imagefiles1)
processed_images = 0

In [8]: import matplotlib.pyplot as plt
from skimage.feature import hog
from skimage import exposure

image_files = image_files[:5]

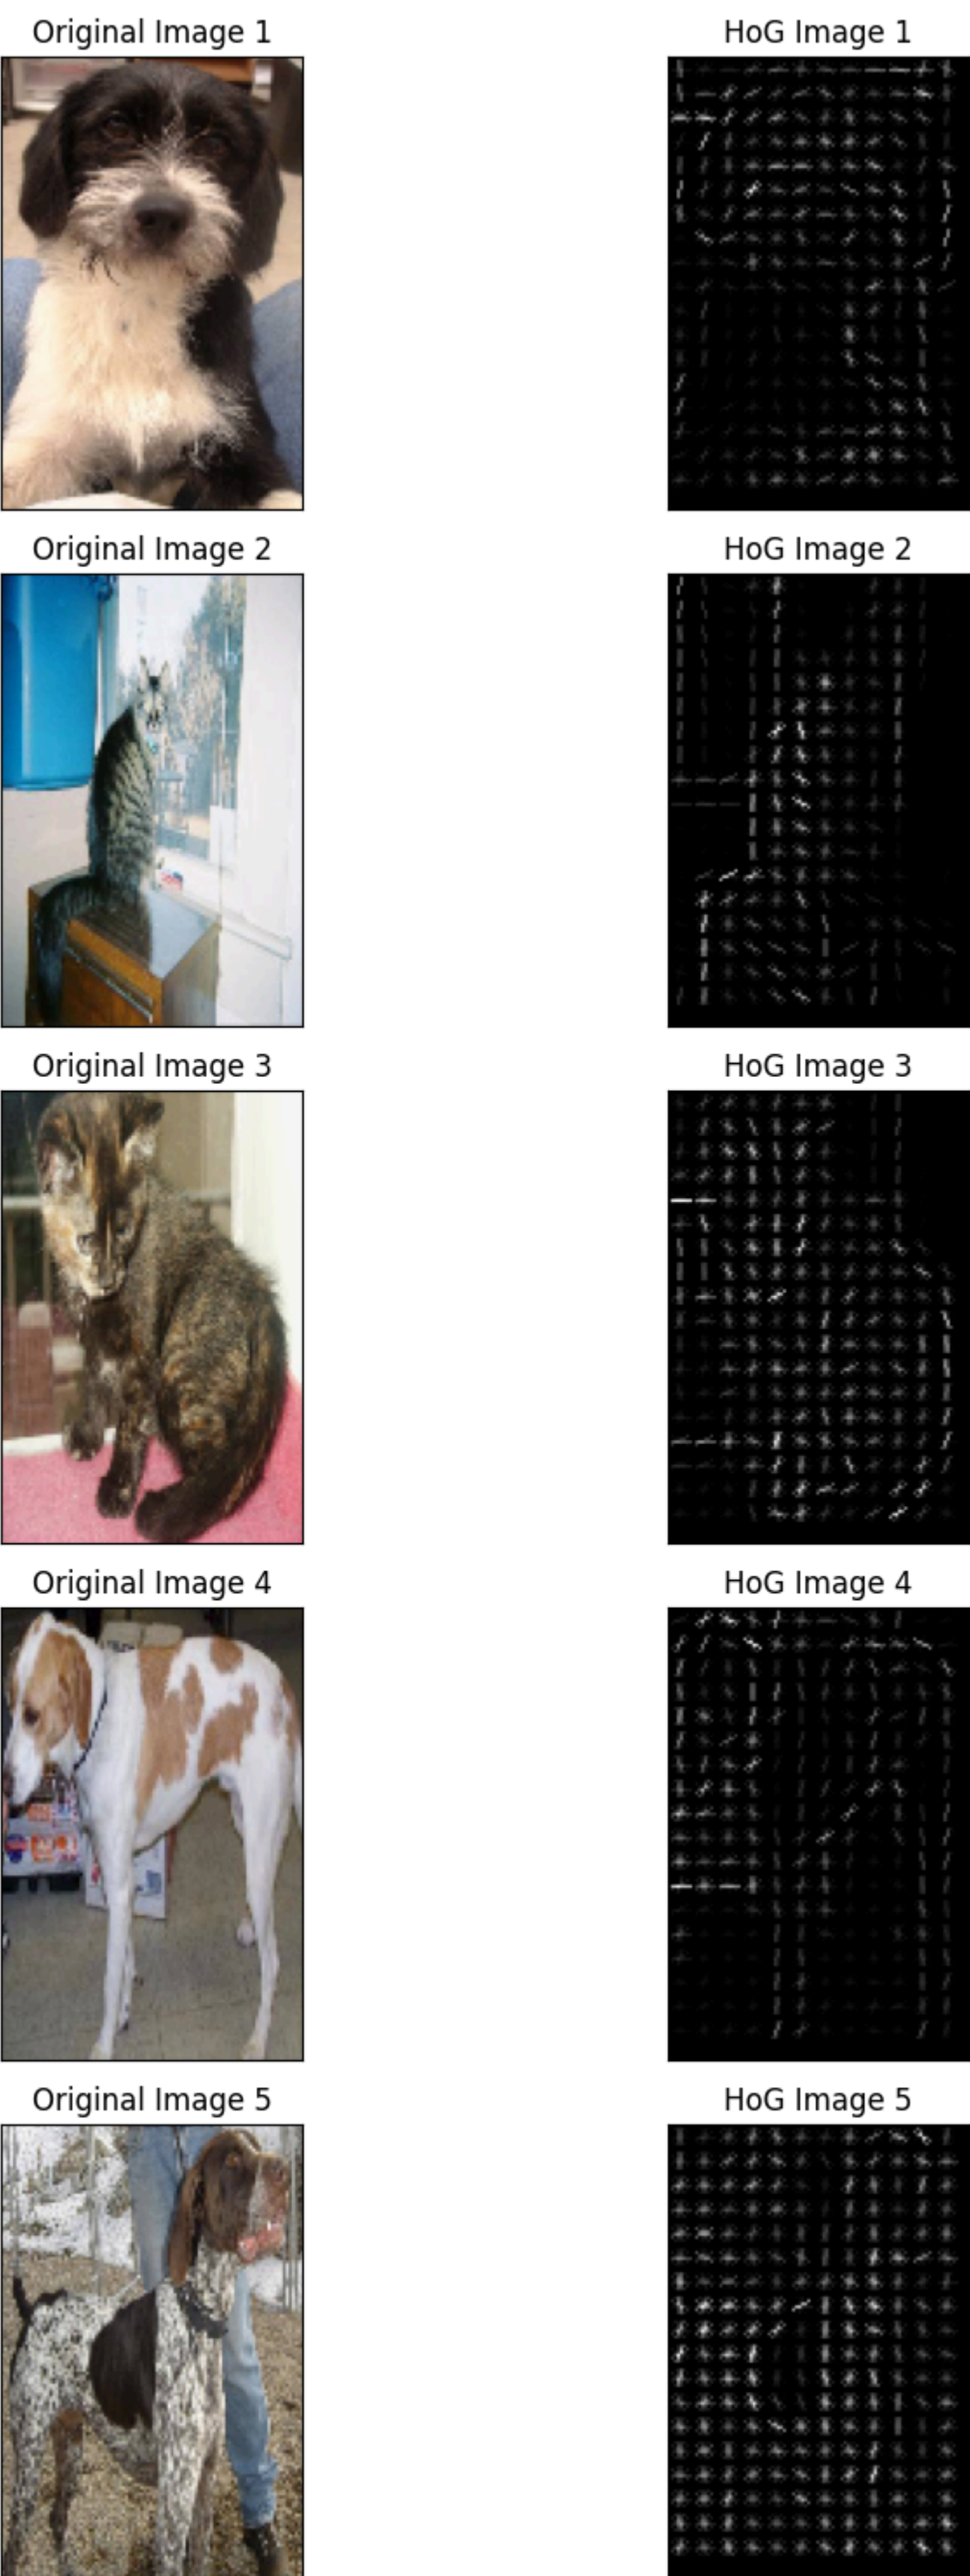
original_images = []
hog_images = []

for image_path in image_files:
    image = cv2.imread(image_path) # Đọc hình ảnh bằng OpenCV
    image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB) # Chuyển đổi sang định dạng RGB
    image = cv2.resize(image, (100,150))
    features, hog_image = compute_hog_feature(image) # Tính toán các đặc trưng HoG
    original_images.append(image)
    hog_images.append(hog_image)

fig, axes = plt.subplots(nrows=5, ncols=2, figsize=(10, 15),
                        subplot_kw={'xticks': [], 'yticks': []})

for i in range(5):
    axes[i, 0].imshow(original_images[i])
    axes[i, 0].set_title(f'Original Image {i+1}')
    axes[i, 1].imshow(hog_images[i], cmap='gray')
    axes[i, 1].set_title(f'HoG Image {i+1}')

plt.tight_layout()
plt.show()
```



```
In [9]: hog_features_list = []
labels = []
for image_path in imagefiles1:
    processed_images += 1
    print(f"Processing image {processed_images}/{total_images}", end="\r")
    image = cv2.imread(image_path)
    image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
    image = cv2.resize(image, (100,150))
    features = compute_hog_features(image)
    hog_features_list.append(features)

    if 'cat' in os.path.basename(image_path):
        labels.append(0) # 0 cho mèo
    elif 'dog' in os.path.basename(image_path):
        labels.append(1) # 1 cho chó

Processing image 25000/25000
```

```
In [10]: X = np.array(hog_features_list)
y = np.array(labels)

In [11]: print(X.shape)
print(y.shape)

(25000, 6732)
(25000,)

In [12]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

In [13]: from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X_train_normalized = scaler.fit_transform(X_train)
X_test_normalized = scaler.transform(X_test)

pca = PCA(n_components=100) # Bạn có thể chọn số lượng thành phần
X_train_pca = pca.fit_transform(X_train_normalized)
X_test_pca = pca.transform(X_test_normalized)

In [14]: from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV
lr = LogisticRegression(max_iter=10000)
search_space = {
    'C': [0.1, 5, 10, 20]
}
GS = GridSearchCV(
    estimator=lr, # Mô hình Logistic Regression
    param_grid=search_space, # Không gian tìm kiếm siêu tham số
)
GS.fit(X_train_pca, y_train)
```

```
Out[14]: > GridSearchCV
> estimator: LogisticRegression
> LogisticRegression
```

```
In [15]: best_params = GS.best_estimator_
best_params
```

```
Out[15]: > LogisticRegression
LogisticRegression(C=0.1, max_iter=10000)
```

```
In [16]: lo_r = LogisticRegression(C=0.1, max_iter=10000)
lo_r.fit(X_train_pca, y_train)
y_lo_r_pred = lo_r.predict(X_test_pca)
print(classification_report(y_test, y_lo_r_pred))
```

	precision	recall	f1-score	support
0	0.71	0.70	0.70	2435
1	0.72	0.73	0.72	2565
accuracy			0.71	5000
macro avg	0.71	0.71	0.71	5000
weighted avg	0.71	0.71	0.71	5000

```
In [17]: svm_1 = SVC(C=1, kernel='rbf')
svm_1.fit(X_train_pca, y_train)
y_pred = svm_1.predict(X_test_pca)
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.78	0.78	0.78	2435
1	0.79	0.79	0.79	2565
accuracy			0.79	5000
macro avg	0.79	0.79	0.79	5000
weighted avg	0.79	0.79	0.79	5000

```
In [18]: from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier()

param_grid = {
    'n_neighbors': [10],
    'weights': ['uniform', 'distance'],
    'metric': ['euclidean', 'manhattan', 'minkowski']
}

grid_search = GridSearchCV(knn, param_grid, cv=5, n_jobs=-1, verbose=1)
grid_search.fit(X_train_pca, y_train)

print(f"Best parameters found: {grid_search.best_params}")

best_knn = grid_search.best_estimator_
y_pred = best_knn.predict(X_test_pca)

print(classification_report(y_test, y_pred))

Fitting 5 folds for each of 6 candidates, totalling 30 fits
Best parameters found: {'metric': 'euclidean', 'n_neighbors': 10, 'weights': 'uniform'}
```

	precision	recall	f1-score	support
0	0.79	0.51	0.62	2435
1	0.65	0.87	0.75	2565
accuracy			0.70	5000
macro avg	0.72	0.69	0.68	5000
weighted avg	0.72	0.70	0.69	5000

```
In [19]: knn = KNeighborsClassifier()

param_grid = {
    'n_neighbors': list(range(1,100)),
    'weights': ['uniform'],
    'metric': ['euclidean']
}

grid_search = GridSearchCV(knn, param_grid, cv=5, n_jobs=-1, verbose=1)
grid_search.fit(X_train_pca, y_train)

print(f"Best parameters found: {grid_search.best_params}")

best_knn = grid_search.best_estimator_
y_pred = best_knn.predict(X_test_pca)

print(classification_report(y_test, y_pred))

Fitting 5 folds for each of 99 candidates, totalling 495 fits
Best parameters found: {'metric': 'euclidean', 'n_neighbors': 4, 'weights': 'uniform'}
```

	precision	recall	f1-score	support
0	0.72	0.61	0.66	2435
1	0.68	0.77	0.72	2565

accuracy			0.69	5800
macro avg	0.70	0.69	0.69	5800
weighted avg	0.69	0.69	0.69	5800